

Big Data and online streaming machine learning

Big Data, Fast Data, All Data

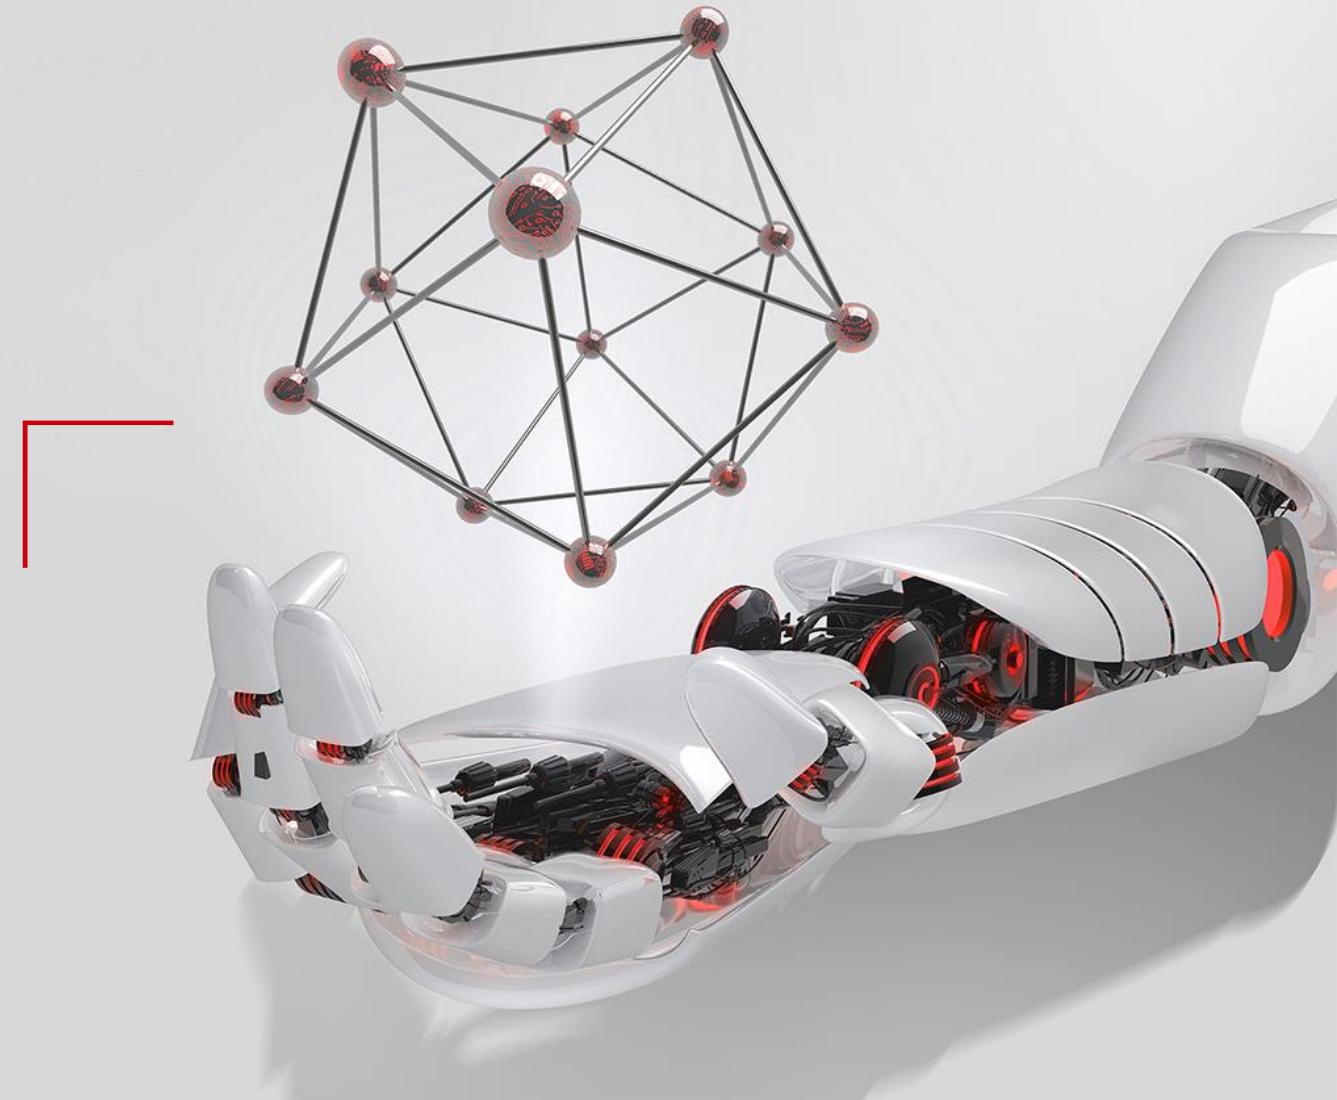
Dr. Cristian Axenie

Senior Research Engineer AI and ML

AI and ML in Big Data Team

IT Software Infrastructure Dept.

Security Level: Open



Introducing the speaker



TUM PhD in
Neuroscience and Robotics,
Summa cum Laude

Specialized in designing and
implementing
AI and ML system for
real-world problems

Academic Research

Head of Research Lab
AI and VR
As of 2017



Lecturer
As of 2017



Postdoctoral Fellow,
Lecturer
2016-2017



Research Assistant (PhD)
2011-2016



Industry Research

Senior Research Engineer
AI, ML & Big Data
As of 2017



Software Engineer
Automotive
2009-2011




Software Engineer
Automotive
2009-2011



Software Engineer
Embedded Systems
2007-2008




(Dis)ambiguating Big Data




Big Data in a Nutshell



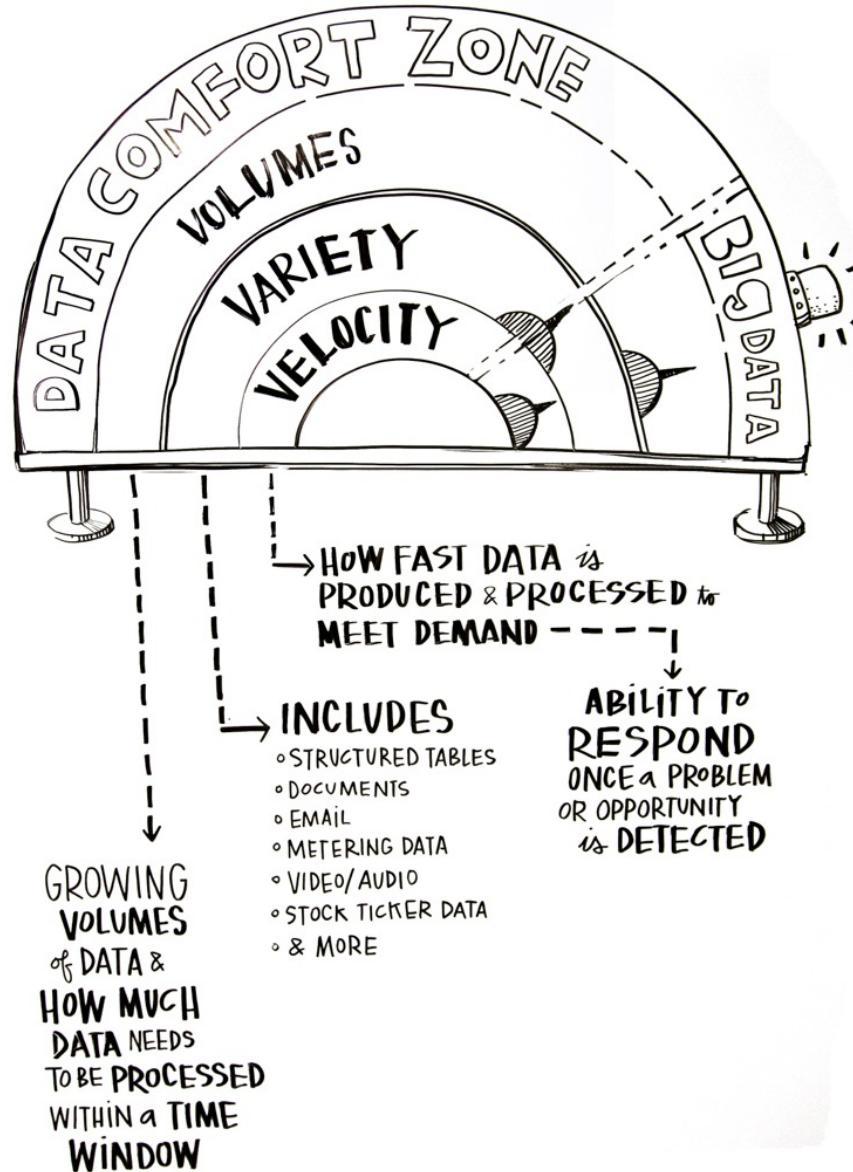
Big Data in a Nutshell



Big Data in a Nutshell



Big Data in a Nutshell



Big Data Stream Processing: A gentle introduction



Big Data Stream Processing: A gentle introduction


Stream processing paradigm simplifies parallel software and hardware by restricting the parallel computation that can be performed.

Given a sequence of data (**a stream**), a series of operations (**functions**) is applied to each element in the stream, in a declarative way, we specify **what** we want to achieve and **not how**.

Big Data Stream Processing: A gentle introduction

Stream processing paradigm simplifies parallel software and hardware by restricting the parallel computation that can be performed.

Given a sequence of data (**a stream**), a series of operations (**functions**) is applied to each element in the stream, in a declarative way, we specify **what** we want to achieve and **not how**.



Big Data Stream Learning: Why is it different?



Big Data Stream Learning: Why is it different?

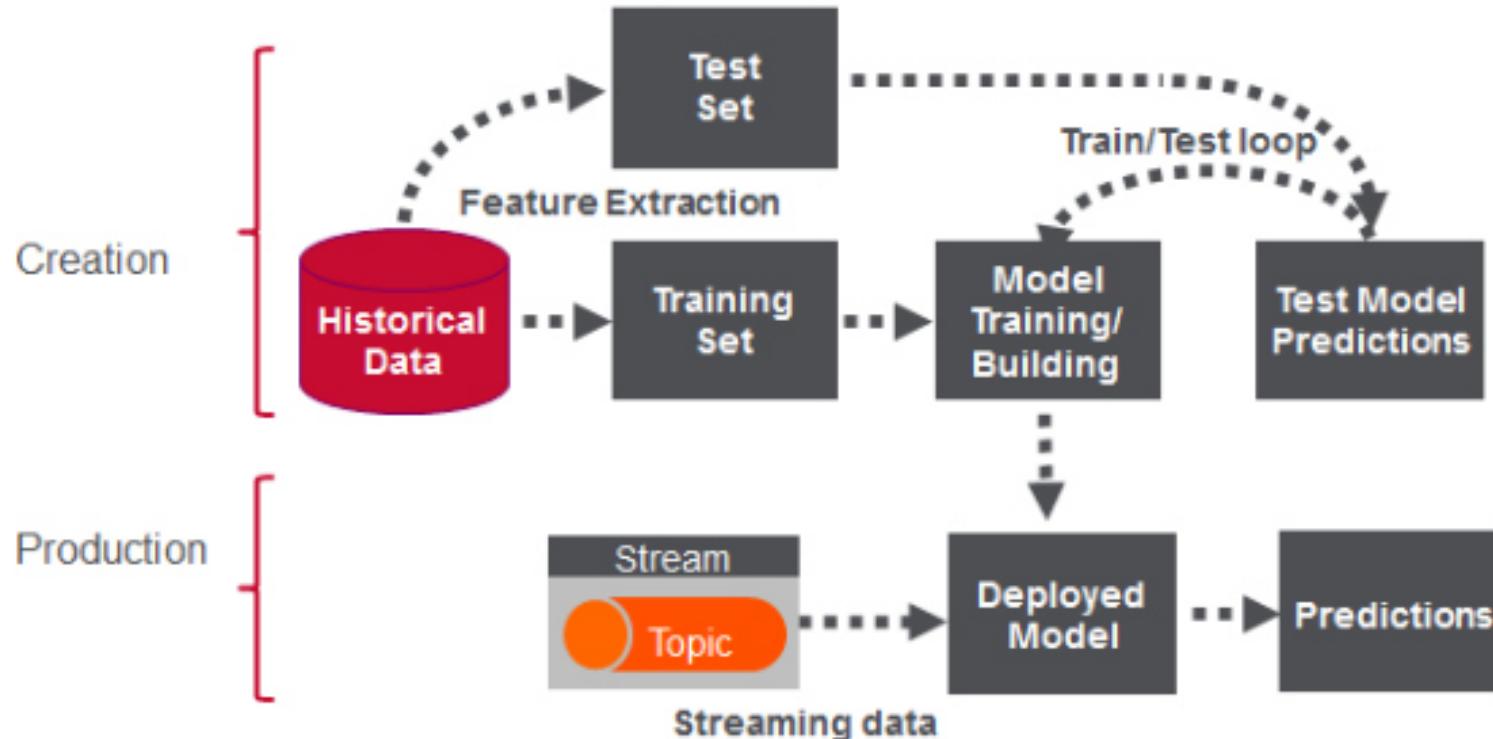
Big Data Stream Learning is more challenging than **batch** or **offline learning**, since the data may **not preserve the same distribution** over the lifetime of the stream.

Moreover, each example coming in a stream can only be **processed once**, or needs to be summarized with a **small memory footprint**, and the learning algorithms must be **efficient**.

Big Data Stream Learning: Why is it different?

Big Data Stream Learning is more challenging than **batch** or **offline learning**, since the data may **not preserve the same distribution** over the lifetime of the stream.

Moreover, each example coming in a stream can only be **processed once**, or needs to be summarized with a **small memory footprint**, and the learning algorithms must be **efficient**.



Big Data Stream Learning: Where's the catch?



Big Data Stream Learning: Where's the catch?

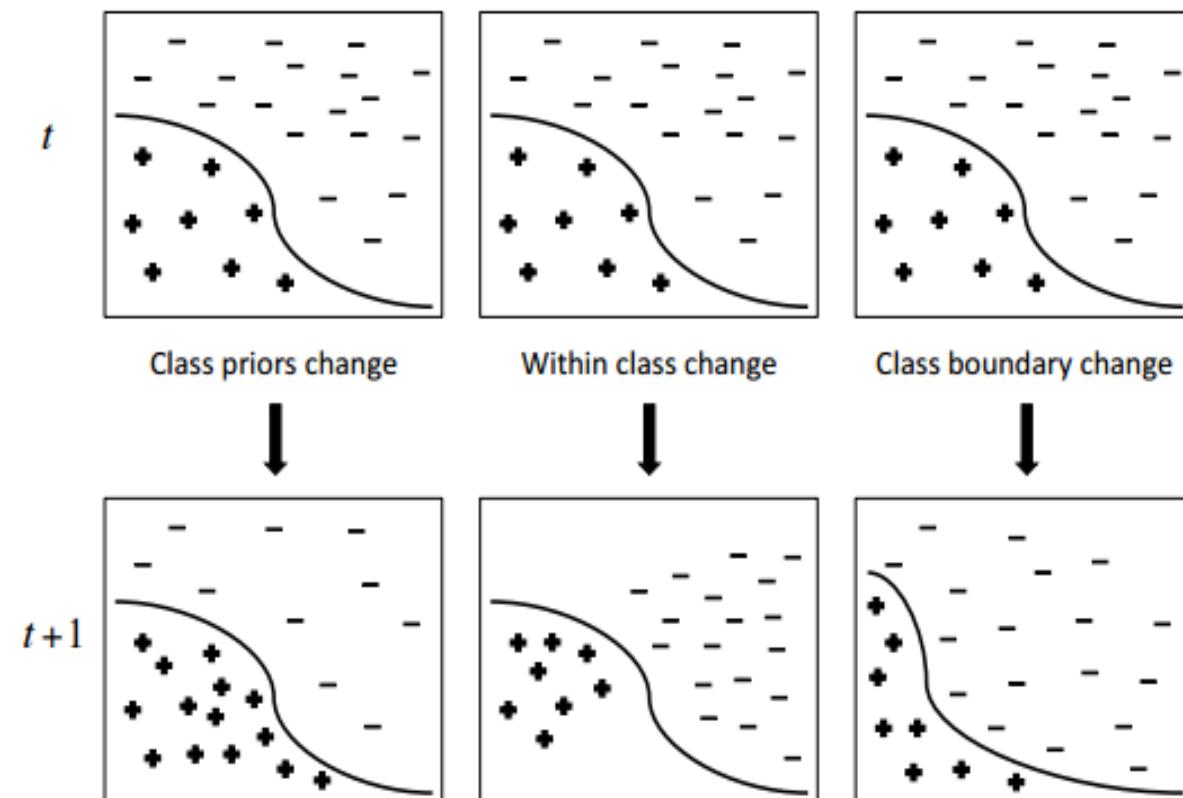
In order to deal with **evolving data streams**, the model learnt from the streaming data must **capture up-to-date trends and transient patterns** in the stream.

Updating the model by incorporating **new examples**, we must also **eliminate the effects of outdated examples** representing outdated concepts through **one-pass**.

Big Data Stream Learning: Where's the catch?

In order to deal with **evolving data streams**, the model learnt from the streaming data must **capture up-to-date trends and transient patterns** in the stream.

Updating the model by incorporating **new examples**, we must also **eliminate the effects of outdated examples** representing outdated concepts through **one-pass**.



Streaming Machine Learning Basics

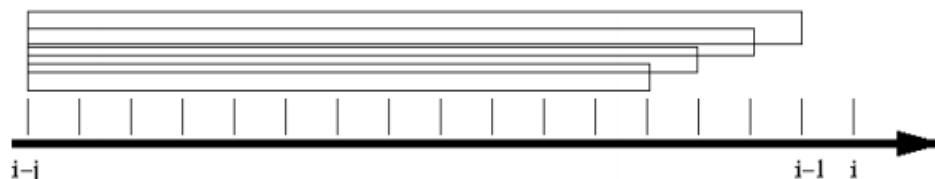
Elements of Streaming Machine Learning

Elements of Streaming Machine Learning


Most strategies use variations of the **sliding window technique**: a window is maintained that keeps the most **recently read examples**, and from which **older examples are dropped** according to some **set of rules**.

Elements of Streaming Machine Learning


Most strategies use variations of the **sliding window technique**: a window is maintained that keeps the most **recently read examples**, and from which **older examples are dropped** according to some **set of rules**.



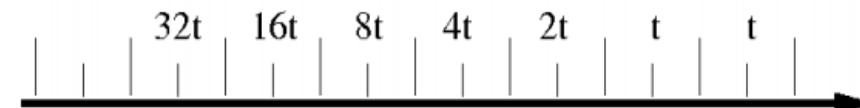
(a) Landmark Window



(b) Sliding Window



(a) Natural Tilted Time Window



b) Logarithmic Tilted Time Window

Elements of Streaming Machine Learning

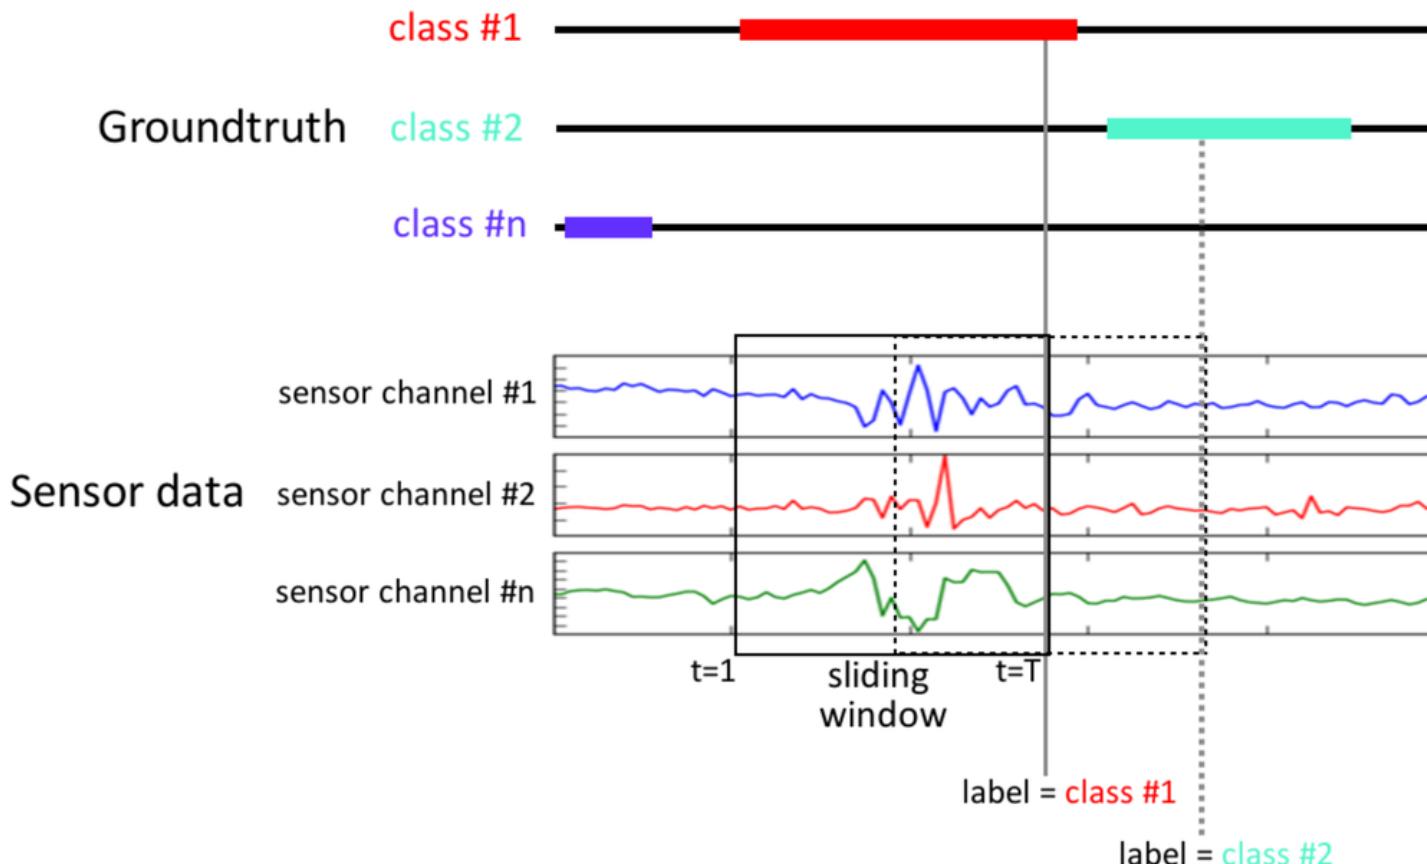
The contents of the **sliding window** can be used for the three tasks:

- 1) to **detect change** (e.g., by using some statistical test on different sub-windows),
- 2) to obtain **updated statistics / criteria** from the **recent examples**, and
- 3) to have **data to rebuild or update the model after data has changed.**

Elements of Streaming Machine Learning

The contents of the **sliding window** can be used for the three tasks:

- 1) to **detect change** (e.g., by using some statistical test on different sub-windows),
- 2) to obtain **updated statistics / criteria** from the **recent examples**, and
- 3) to have **data to rebuild or update the model** after data has changed.



Elements of Streaming Machine Learning



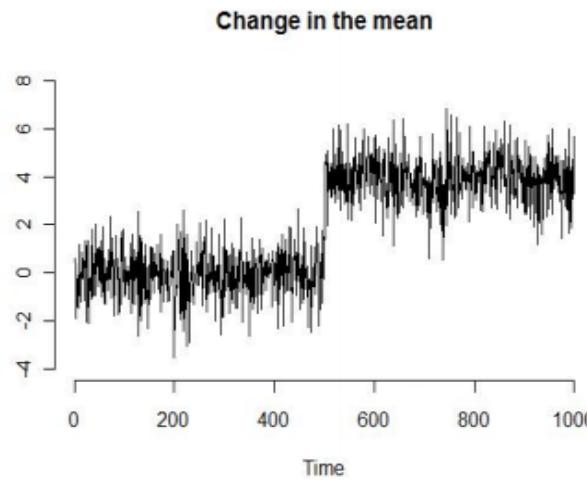
Normally, the user is caught in a **tradeoff without solution**:

- a **small size** (so that the **window reflects** accurately the current **distribution**)
- a **large size** (so that many **examples** are available to work on, **increasing accuracy** in periods of stability).

Elements of Streaming Machine Learning

Normally, the user is caught in a **tradeoff without solution**:

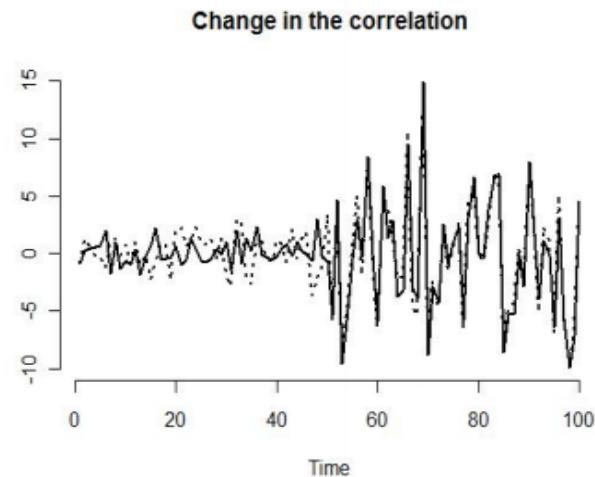
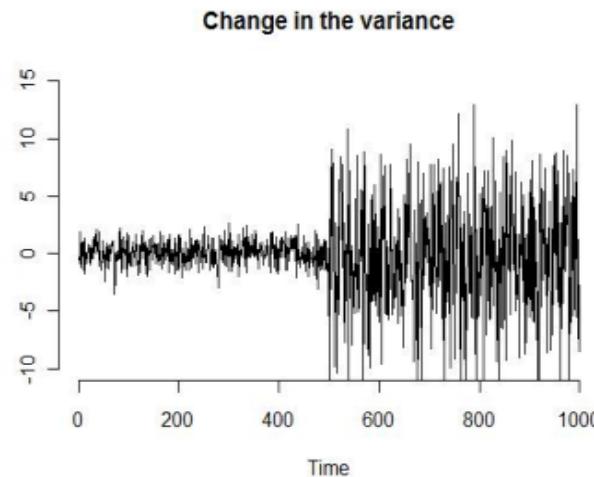
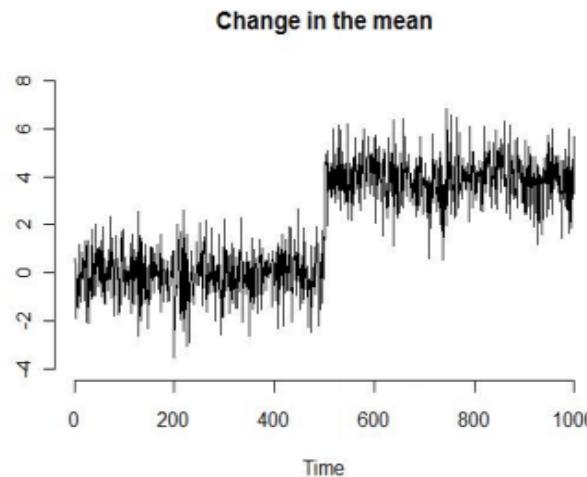
- a **small size** (so that the **window reflects** accurately the current **distribution**)
- a **large size** (so that many **examples** are available to work on, **increasing accuracy** in periods of stability).



Elements of Streaming Machine Learning

Normally, the user is caught in a **tradeoff without solution**:

- a **small size** (so that the **window reflects** accurately the current **distribution**)
- a **large size** (so that many **examples** are available to work on, **increasing accuracy** in periods of stability).



Currently, it has been proposed to use **windows of variable size**.

Streaming Machine Learning in Action



Streaming Machine Learning I

Example use-cases **Pollution GO**



Urban Observatory project at Newcastle University:

- request sensory data to tackle specific challenges, such as flooding or air quality within **Newcastle Metropolitan area:**

<http://uoweb1.ncl.ac.uk/>

- assuming we interrogate **two sites** (locations) in the **Newcastle Urban Observatory**, marked as **red hexagons** in the images.

You can see what is the **available sensory information** at each site.



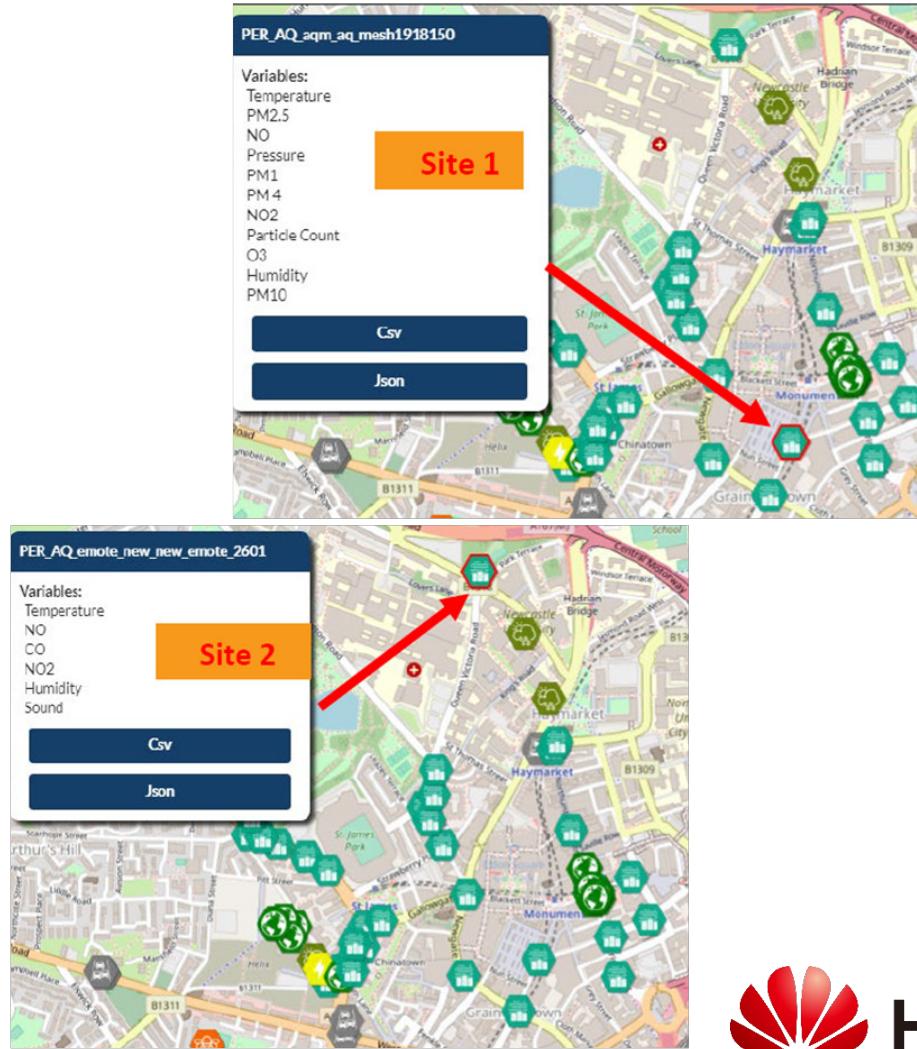
Streaming Machine Learning I

Example use-cases **Pollution GO**

The task is to **learn pair-wise correlations** among the **available sensors**.

The **system performs unsupervised learning of functional relationships** between two **input sensory streams** (e.g. NO in Site1 and Temperature in Site 1).

This **neural network** based system can be employed in **various solutions** as a **tool to learn pair-wise sensory correlations** among sensors within / between spatial locations (e.g. CO in Site 2 and O3 in Site 1).




Streaming Machine Learning I

Example use-cases **Pollution GO**

The **power of the approach** is that one can **learn / extract sensory correlations in various constellations**:

- Learn between location within sensor correlations
 - example NO₂ for **site 1** and **site 2**, corresponding to X and Y respectively, over a range



Streaming Machine Learning II

Example use-cases **Traffic GO**

The approach can learn / extract sensory correlations in various scenarios.

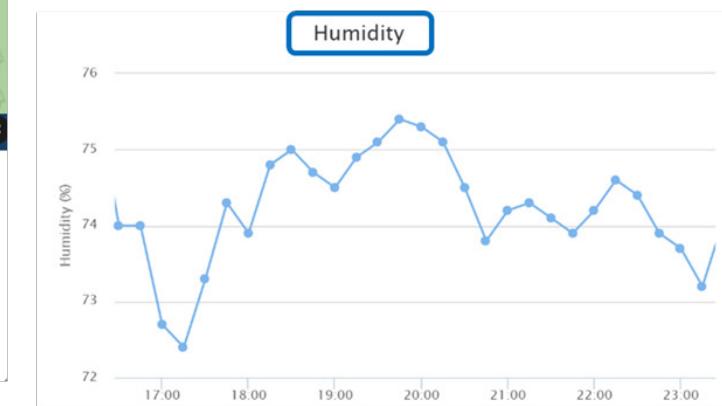
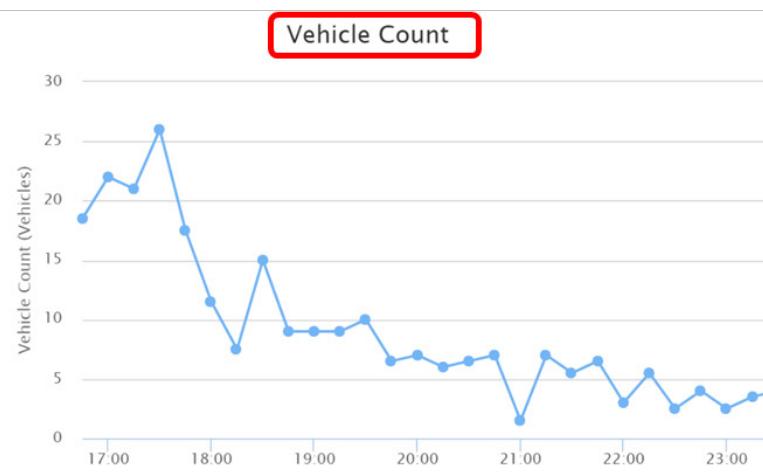
In the **traffic scenario** we propose to **learn the correlation** between **Environment parameters** (NO, O₃, NO₂, NO_x), **Weather** (Humidity, Rain) and the **Traffic Flow** (number of vehicles) at a site.

Once **learnt the correlation** we can use it to **infer traffic flow** in regions where **we do not have traffic sensors** installed but **all other sensors** are present.






Streaming Machine Learning II

Example use-cases **Traffic GO**



Streaming Machine Learning II

Once learnt the correlation we can use it to infer traffic flow in regions where we do not have traffic sensors installed but all other sensors are present (e.g. Humidity and NO₂). 



If the **traffic sensor is failing / defect**,
The system uses previously learnt relations
to infer a plausible prediction.



Streaming Machine Learning III


Example use-case **Elevator Analytics**

The approach can learn / extract sensory correlations for Predictive Maintenance of Elevator Doors.

For an elevator car door the system can learn the correlation between:

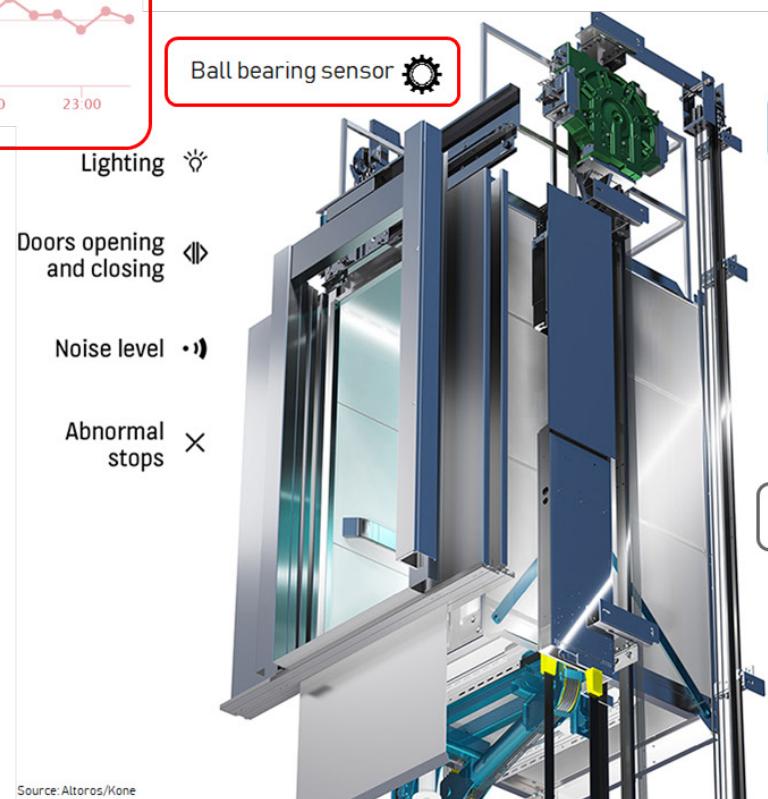

- Electromechanical sensors
(Door Ball Bearing Sensor)
- Ambiance(Humidity)
- Physics(Vibration)

Once learnt the correlation in operational settings we can use it to infer anomalous operation of the doors.





Streaming Machine Learning III

Example use-case **Elevator Analytics**



We have the operation data, timeseries sampled at 10Hz in high-peak and evening elevator usage in a building (between 16:30 and 23:30).

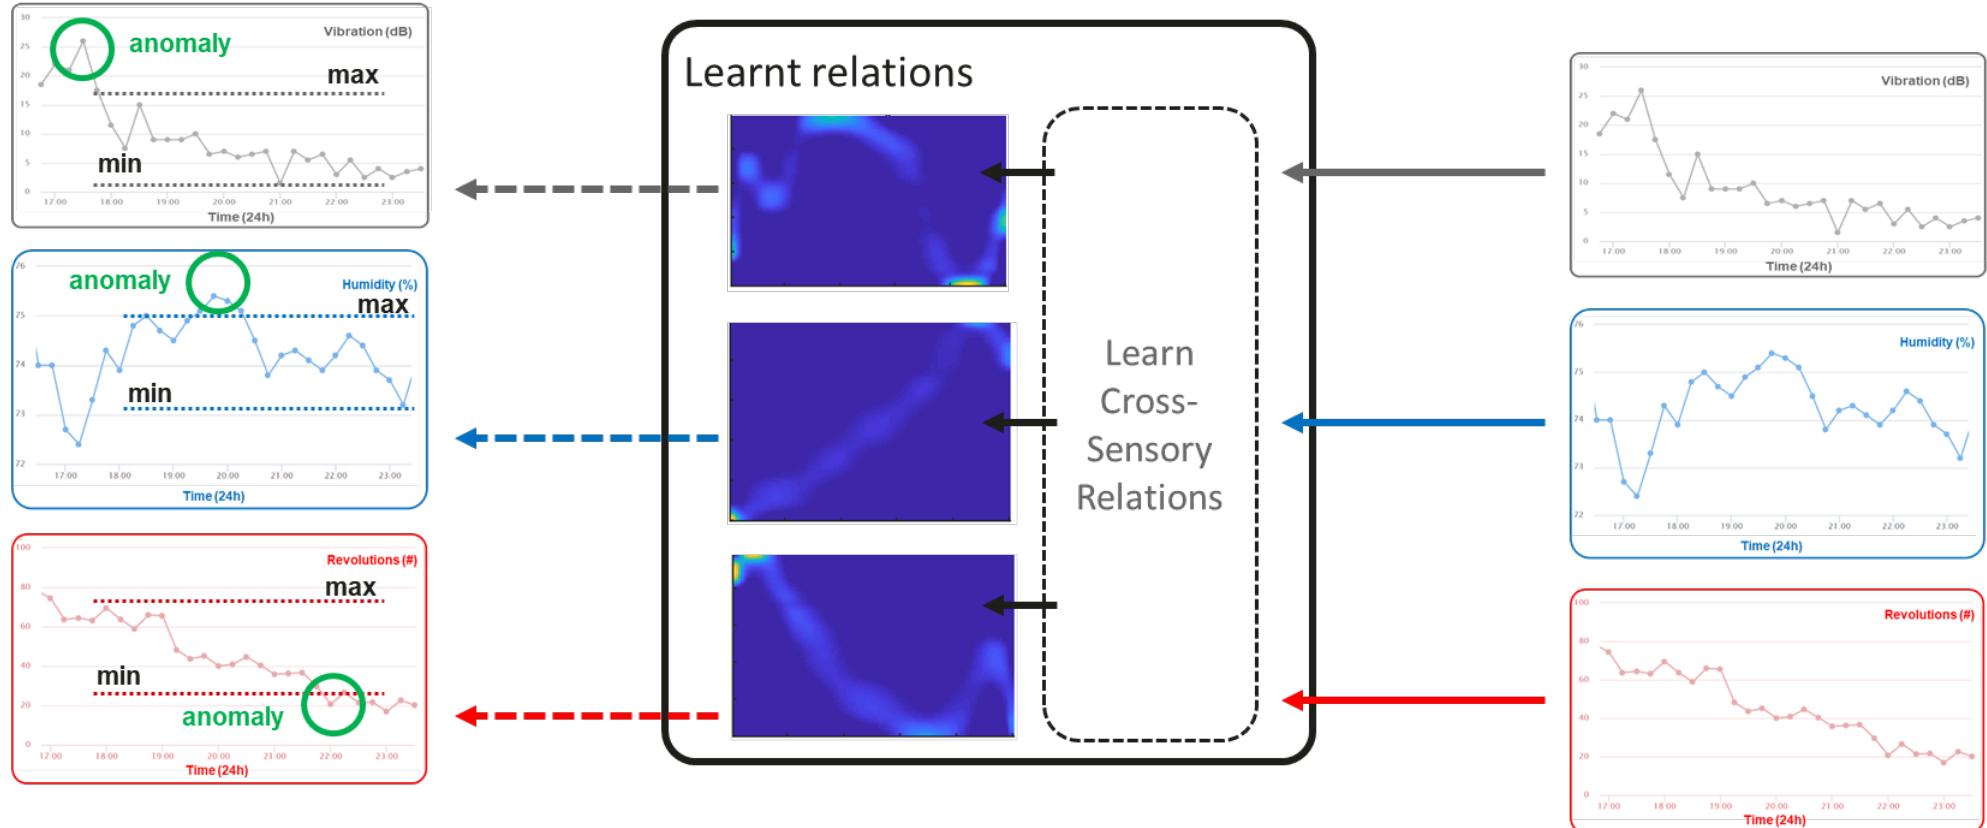
Source: Altoros/Kone



Streaming Machine Learning III

Example use-case **Elevator Analytics**

Once learnt the correlation we can use it to trigger alarms for anomalies / outliers.



Conclusions

Conclusions

Induced by ubiquitous scenarios *finite training sets*, *static models*, and *stationary distributions* must be completely redefined.



Conclusions

Induced by ubiquitous scenarios *finite training sets*, *static models*, and *stationary distributions* must be completely redefined.

The characteristics of the streaming data entail a new vision due to the fact that:

Conclusions

Induced by ubiquitous scenarios *finite training sets*, *static models*, and *stationary distributions* must be completely redefined.

The characteristics of the streaming data entail a new vision due to the fact that:

- Data are made available through **unlimited streams** that **continuously flow**, eventually at **high speed**, over time;

Conclusions

Induced by ubiquitous scenarios *finite training sets*, *static models*, and *stationary distributions* must be completely redefined.

The characteristics of the streaming data entail a new vision due to the fact that:

- Data are made available through **unlimited streams** that **continuously flow**, eventually at **high speed**, over time;
- The **underlying regularities may evolve** over time rather than being stationary;

Conclusions

Induced by ubiquitous scenarios *finite training sets*, *static models*, and *stationary distributions* must be completely redefined.

The characteristics of the streaming data entail a new vision due to the fact that:

- Data are made available through **unlimited streams** that **continuously flow**, eventually at **high speed**, over time;
- The **underlying regularities may evolve** over time rather than being stationary;
- The data can **no longer** be considered as **independent and identically distributed**;

Conclusions

Induced by ubiquitous scenarios *finite training sets*, *static models*, and *stationary distributions* must be completely redefined.

The characteristics of the streaming data entail a new vision due to the fact that:

- Data are made available through **unlimited streams** that **continuously flow**, eventually at **high speed**, over time;
- The **underlying regularities may evolve** over time rather than being stationary;
- The data can **no longer** be considered as **independent and identically distributed**;
- The **data** are now often **spatially** as well as **time situated**.

Conclusions

Induced by ubiquitous scenarios *finite training sets*, *static models*, and *stationary distributions* must be completely redefined.

The characteristics of the streaming data entail a new vision due to the fact that:

- Data are made available through **unlimited streams** that **continuously flow**, eventually at **high speed**, over time;
- The **underlying regularities may evolve** over time rather than being stationary;
- The data can **no longer** be considered as **independent and identically distributed**;
- The **data** are now often **spatially** as well as **time situated**.

A new era of machine learning?

Thank you.

Bring digital to every person, home, and organization for a fully connected, intelligent world.


**Copyright©2018 Huawei Technologies Co., Ltd.
All Rights Reserved.**

The information in this document may contain predictive statements including, without limitation, statements regarding the future financial and operating results, future product portfolio, new technology, etc. There are a number of factors that could cause actual results and developments to differ materially from those expressed or implied in the predictive statements. Therefore, such information is provided for reference purpose only and constitutes neither an offer nor an acceptance. Huawei may change the information at any time without notice.




Stream Processing Engines

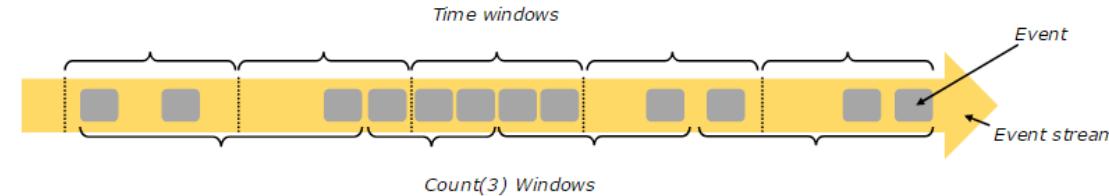
Stream Processing Engines



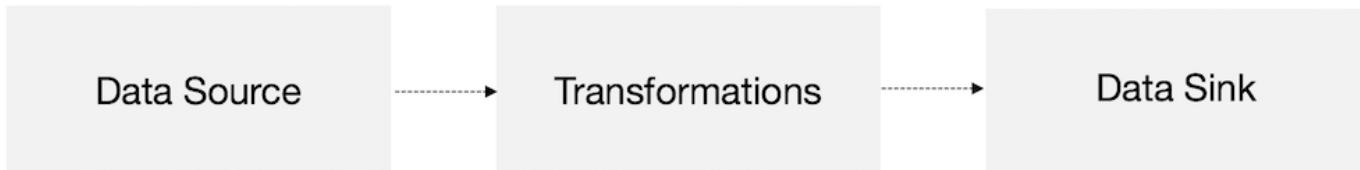
Stream Processing Engines



Aggregating events (e.g., counts, sums) works differently on streams because it is **impossible to count all** (unbounded).

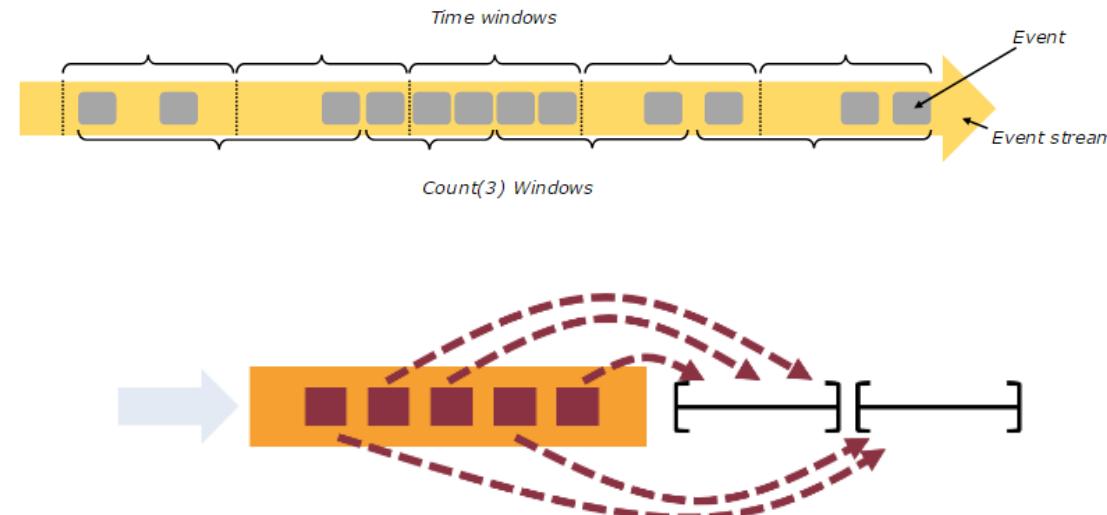


Stream Processing Engines

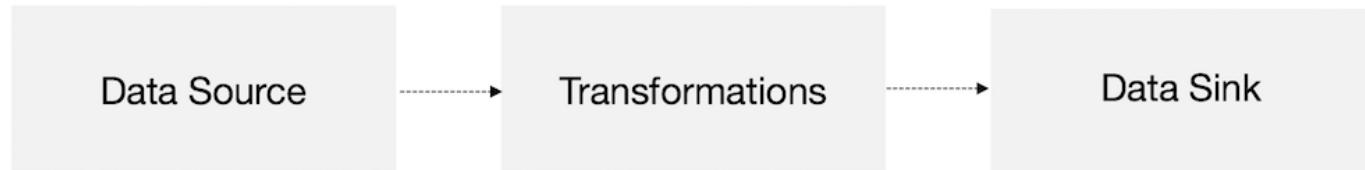


Aggregating events (e.g., counts, sums) works differently on streams because it is **impossible to count all** (unbounded).

Stream processing and **windowing** makes it easy to compute accurate results over streams where **events arrive out of order** and where **events may arrive delayed**.



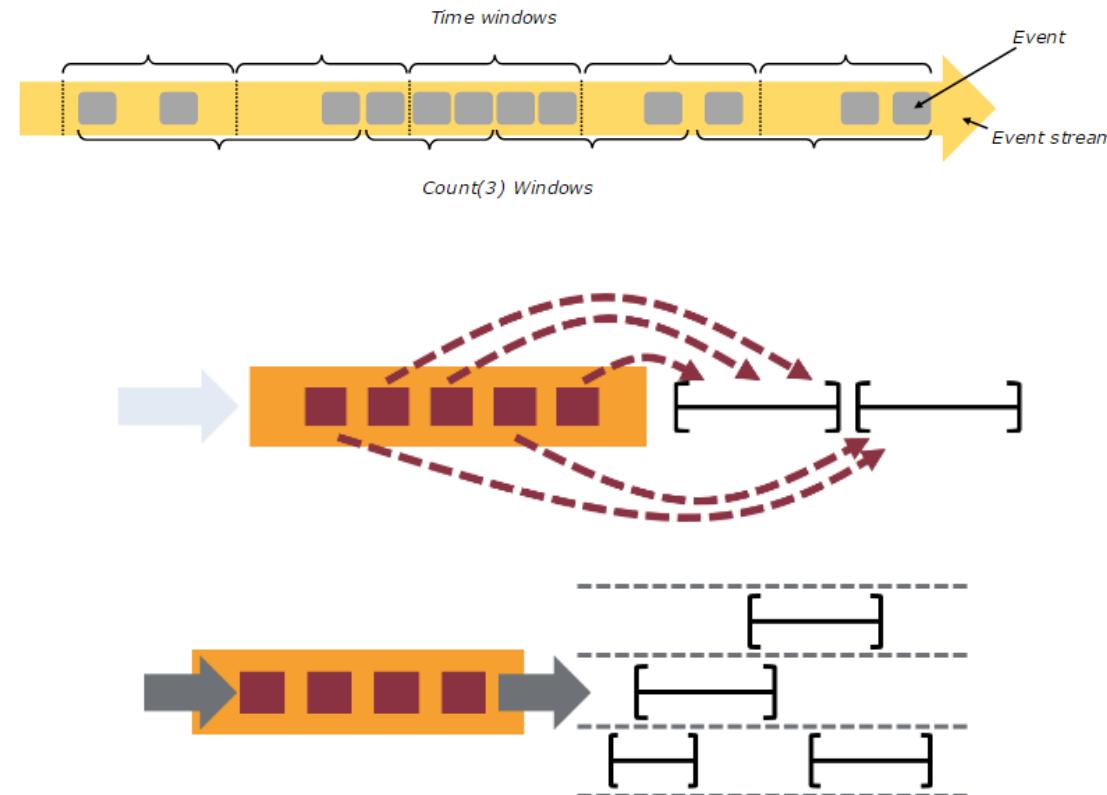
Stream Processing Engines



Aggregating events (e.g., counts, sums) works differently on streams because it is **impossible to count all** (unbounded).

Stream processing and **windowing** makes it easy to compute accurate results over streams where **events arrive out of order** and where **events may arrive delayed**.

Windowing based on time, count, and data-driven windows. Windows can be customized with **flexible triggering** conditions to support **sophisticated streaming patterns**.



Stream Processing Engines



When executed, Flink programs are mapped to **streaming dataflows**, consisting of **streams** and **transformation operators**.

Each **dataflow** starts with one or more **sources** and ends in one or more **sinks**.

The dataflows resemble arbitrary directed acyclic graphs (DAGs).

Stream Processing Engines



When executed, Flink programs are mapped to **streaming dataflows**, consisting of **streams** and **transformation operators**.

Each **dataflow** starts with one or more **sources** and ends in one or more **sinks**.

The dataflows resemble arbitrary directed acyclic graphs (DAGs).

```
DataStream<String> lines = env.addSource(  
    new FlinkKafkaConsumer<>(...)); } Source  
  
DataStream<Event> events = lines.map((line) -> parse(line)); } Transformation  
  
DataStream<Statistics> stats = events  
    .keyBy("id")  
    .timeWindow(Time.seconds(10))  
    .apply(new MyWindowAggregationFunction()); } Transformation  
  
stats.addSink(new RollingSink(path)); } Sink
```

Stream Processing Engines




When executed, Flink programs are mapped to **streaming dataflows**, consisting of **streams** and **transformation operators**.

Each **dataflow** starts with one or more **sources** and ends in one or more **sinks**.

The dataflows resemble arbitrary directed acyclic graphs (DAGs).

```
DataStream<String> lines = env.addSource(  
    new FlinkKafkaConsumer<>(...)); } Source  
  
DataStream<Event> events = lines.map((line) -> parse(line)); } Transformation  
  
DataStream<Statistics> stats = events  
    .keyBy("id")  
    .timeWindow(Time.seconds(10))  
    .apply(new MyWindowAggregationFunction()); } Transformation  
  
stats.addSink(new RollingSink(path)); } Sink
```



Stream Processing Engines

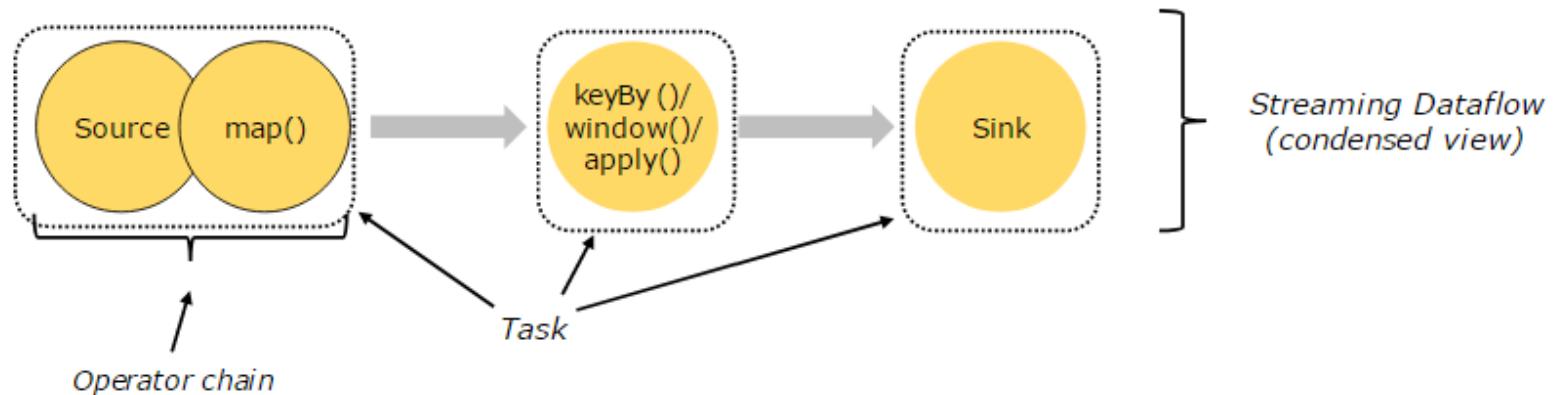


For distributed execution, Flink chains operator **subtasks** together into **tasks**. Each task is executed by one **thread**.

Stream Processing Engines



For distributed execution, Flink chains operator **subtasks** together into **tasks**. Each task is executed by one **thread**.

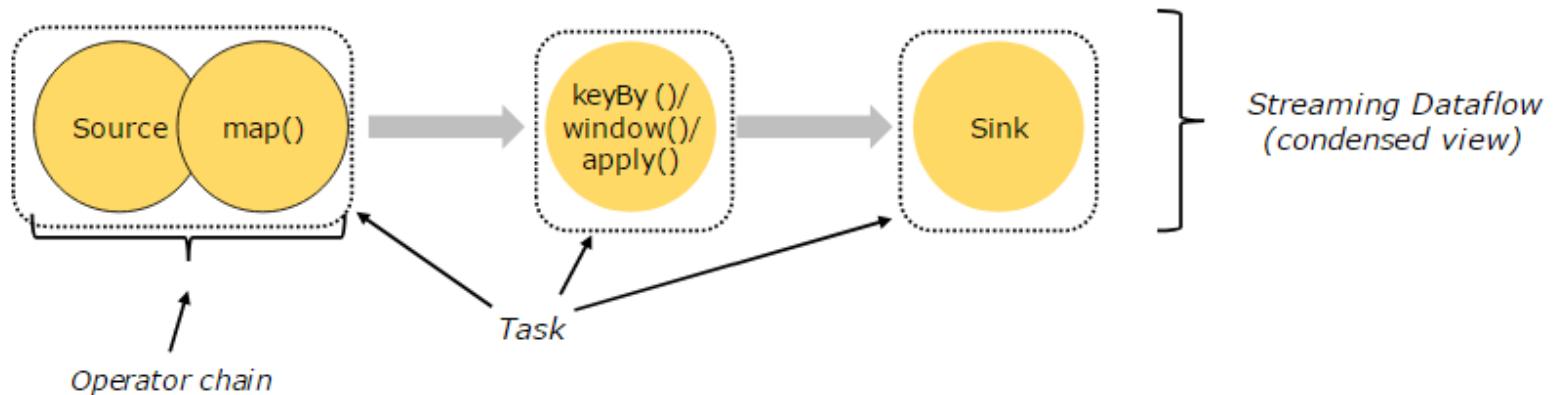


Stream Processing Engines



For distributed execution, Flink chains operator **subtasks** together into **tasks**. Each task is executed by one **thread**.

Chaining operators together into tasks is a useful optimization: it **reduces** the **overhead** of thread-to-thread handover and **buffering**, and **increases overall throughput** while **decreasing latency**.



Stream Processing Engines



For distributed execution, Flink chains operator **subtasks** together into **tasks**. Each task is executed by one **thread**.

Chaining operators together into tasks is a useful optimization: it **reduces** the **overhead** of thread-to-thread handover and **buffering**, and **increases overall throughput** while **decreasing latency**.

